
Sequence Containers

Reid Watson
(rawatson@stanford.edu)

Administrivia

- Assignment one is out today!
- Quick Demo
- Due: Wednesday October 30th, 11:59 PM

Administrivia

- LalR help available:
 - Sunday 8-12
 - Monday 8-10
 - More times to be announced, check the website for details
 - Email help available
 - cs106l-staff-spr2013@lists.stanford.edu
 - rawatson@stanford.edu
 - Email cs106l-staff for best response time with assignment help
-

Review: Sequence Containers

- A container class allows you to store an arbitrary number of things
 - A sequence container is a container whose elements can be accessed sequentially.
 - Sequence containers include vectors, stacks, queues, lists, and priority queues (among others).
-

What I Want To Show You

- Why the Stanford library exists
 - How to use STL sequence containers instead of the Stanford Library
 - We'll look at the differences between STL/Stanford using stack and vector, and we'll also examine a new STL class, deque
 - The performance characteristics of various sequence containers, and why you might choose one over another
-

Why the Stanford Library Exists

Students often ask:

“Why do we need to use the Stanford libraries
in CS106B/X?”

Why the Stanford Library Exists

- The Stanford libraries include things not found in the STL (`Grid`, `getInteger` and friends, graphics).
 - Many parts of the Stanford library give up performance for simplicity
 - Debugging Stanford library code can be much easier than debugging STL code (see `HugeError.pro`)
-

Container #1: Stack

First, let's talk about how stacks are represented in the STL.

STL <stack>: What's Similar

What you want to do	Stanford Stack<int>	STL stack<int>
Create a stack	<code>Stack<int> x;</code>	<code>stack<int> x;</code>
Get the size of a stack	<code>int size = x.size();</code>	<code>int size = x.size();</code>
Check if a stack is empty	<code>if (x.isEmpty()) ...</code>	<code>if (x.empty()) ...</code>
Push a value on the stack	<code>x.push(42);</code>	<code>x.push(42);</code>
Peek at the top element without popping it	<code>int top = x.peek();</code>	<code>int top = x.top();</code>
Pop off the top element and ignore its value	<code>x.pop();</code>	<code>x.pop();</code>

STL <stack>: What's Different

What you want to do	Stanford Stack<int>	STL stack<int>
Clear the stack	<code>x.clear();</code>	<code>while(!x.empty()) x.pop();</code>
Convert the stack to a string	<code>string s = x. toString();</code>	<code>string s; while(!x.empty()) { s += x.top(); s += " "; x.pop(); }</code>
Pop and save the value	<code>int top = x.pop();</code>	<code>int top = x.top(); x.pop();</code>

STL <stack>: Usage

Let's look at a quick demo in STLStack.pro

STL <stack>: Why the differences?

Looking at the differences between the STL and the Stanford libraries can help you understand the guiding principles behind how each of these libraries were designed.

STL <stack>: Why the differences?

Why is there no .clear() function for stacks?

STL <stack>: Why the differences?

Why is there no .clear() function for stacks?

- Conceptually, clearing isn't part of the interface to a stack
- It's very easy to write your own clear function:

```
// stack<int> s = ...;  
while (!s.empty()) {  
    s.pop();  
}
```

STL <stack>: Why the differences?

Why doesn't pop return the value it removed?

STL <stack>: Why the differences?

Why doesn't pop return the value it removed?

- The caller might not need the value, in which case returning the value would be wasteful.
- It's easy to write code which pops and saves the value.

```
// stack<int> s = ...;  
int value = s.top();  
s.pop();
```

STL <stack>: Why the differences?

Why isn't there a toString function?

STL `<stack>`: Why the differences?

Why isn't there a `toString` function?

- Implementing `toString` would require that the type stored in the stack could be converted to a string
 - For example, you can convert a `stack<int>` to a string because you can convert an `int` to a string.
 - It's tough to say what the "proper" way to convert a stack to a string is
-

Container #2: Vector

First, let's talk about how vectors are represented in the STL.

STL <vector>: What's Similar

What you want to do	Stanford Vector<int>	STL vector<int>
Create an empty vector	<code>Vector<int> v;</code>	<code>vector<int> v;</code>
Create a vector with n copies of zero	<code>Vector<int> v(n);</code>	<code>vector<int> v(n);</code>
Create a vector with n copies of a value k	<code>Vector<int> v(n, k);</code>	<code>vector<int> v(n, k);</code>
Add a value k to the end of the vector	<code>v.add(k);</code>	<code>v.push_back(k);</code>
Clear a vector	<code>v.clear();</code>	<code>v.clear();</code>
Get the element at index i (verify that i is in bounds)	<code>int k = v.get(i);</code> <code>int k = v[i];</code>	<code>int k = v.at(i);</code>
Check if the vector is empty	<code>if (v.isEmpty()) ...</code>	<code>if (v.empty()) ...</code>
Replace the element at index i (verify that i is in bounds)	<code>v.get(i) = k;</code> <code>v[i] = k;</code>	<code>v.at(i) = k;</code>

STL `<vector>`: What's Different

Get the element at index <i>i</i> without bounds checking	<code>// Impossible!</code>	<code>int a = x[i];</code>
Change the element at index <i>i</i> without bounds checking	<code>// Impossible!</code>	<code>x[i] = v;</code>
Apply a function to each element in <i>x</i>	<code>x.mapAll(fn)</code>	<code>// We'll talk about this in another lecture...</code>
Concatenate vectors <i>v1</i> and <i>v2</i>	<code>v1 += v2;</code>	<code>// We'll talk about this in another lecture...</code>
Add an element to the beginning of a vector	<code>// Impossible!</code>	<code>// Impossible!</code>

STL <vector>: Usage

Let's look at a quick demo in `STLVector.pro`

STL `<vector>`: Why the differences?

Why doesn't vector have bounds checking?

STL `<vector>`: Why the differences?

Why doesn't vector have bounds checking?

- If you write your program correctly, bounds checking will do nothing but make your code run slower

STL `<vector>`: Why the differences?

Why is there no `push_front` method?

STL `<vector>`: Why the differences?

Why is there no `push_front` method?

- This is a bit more complicated

The Mystery of push_front


Pushing an element to the front of the vector requires shifting all other elements in the vector down by one, which can be **very** slow

To demonstrate this, let's say we had this nice little vector:

6	7	5	3	0	9
---	---	---	---	---	---


The Mystery of push_front

Now, let's say that `push_front` existed, and that you wanted to insert an 8 at the beginning of this vector.




The Mystery of push_front

First, we may have to expand the capacity of the vector



The Mystery of push_front

Then, we'll need to shift every single element down one position



The Mystery of push_front

Finally, we can actually insert the element we wanted to insert.

```
v.push_front(8).
```



Just how bad is push_front?

```
// Adding to the back  
for (int i = 0; i < N; i++)  
    v.push_back(i);
```

```
// Or: Adding to the front  
for (int i = 0; i < N; i++)  
    v.insert(v.begin(), i);
```

```
// How big can the difference be?
```

Just how bad is push_front?

	push_front	push_back
N = 1000	0.01	0
N = 10000	0.89	0.01
N = 100000	117.98	0.04
N = 1000000	Hours	0.31
N = 10000000	Years	3.16

You can see the difference between an $O(n^2)$ algorithm and an $O(n)$ algorithm!

STL <deque>: What's a deque?


- A deque (pronounced "deck") is a **double ended queue**
 - Unlike a vector, it's possible (and fast) to `push_front`
 - The implementation of a deque isn't as straightforward as a vector though
-

STL <deque>: Usage

Let's look at a quick demo in `STLDeque.cpp`


STL <deque>: Implementation

There's no single specification for representing a deque, but it might be laid out something like this




STL <deque>: Implementation

You could support efficient insertion by keeping some reserved space in front of the vector representing the first elements of the deque



STL <deque>: Implementation

You could support efficient insertion by keeping some reserved space in front of the vector representing the first elements of the deque



STL <deque>: Performance

- We can now use the `push_front` function, and it will run much faster than if we had used a vector.
 - However, if all you're doing is iterating, resizing, and `push_backing`, then using a vector will be faster.
 - Let's see how this looks in real world performance numbers
-

push_front: vector and deque

```
// Vector test code
vector<int> v;
// Insert at the start of the vector
for (int i = 0; i < N; i++)
    v.insert(v.begin(), i);
// Clear by using pop_front (erase)
for (int i = 0; i < N; i++)
    v.erase(v.begin());
```

push_front: vector and deque

```
// Deque test code
deque<int> d;
// Insert elements using push_front
for (int i = 0; i < N; i++)
    d.push_front(i);
// Clear by using pop_front
for (int i = 0; i < N; i++)
    d.pop_front();
```

push_front: vector and deque

	<vector>	<deque>
N = 1000	0.02	0
N = 10000	2.12	0.01
N = 100000	264.9	0.04
N = 1000000	Years	0.44
N = 10000000	Millenia	5.54

Element Access: vector and deque

```
vector<int> v;
```

```
deque<int> d;
```

```
for (int i = 0; i < N; i++)
```

```
    v[i] = i;
```

```
for (int i = 0; i < N; i++)
```

```
    d[i] = i;
```

Access: vector and deque

	<vector>	<deque>
N = 1000	0.02	0.14
N = 10000	0.28	1.32
N = 100000	3.02	13.22
N = 1000000	30.84	133.30

push_back: vector and deque

```
// Vector test code
vector<int> v;
// Insert elements using push_back
for (int i = 0; i < N; i++)
    v.push_back(i);
// Clear by using pop_back
for (int i = 0; i < N; i++)
    v.pop_back();
```

push_back: vector and deque

```
// Deque test code
deque<int> d;
// Insert elements using push_back
for (int i = 0; i < N; i++)
    d.push_back(i);
// Clear by using pop_back
for (int i = 0; i < N; i++)
    d.pop_back();
```

push_back: vector and deque

	<vector>	<deque>
N = 1000	0.02	0.02
N = 10000	0.20	0.20
N = 100000	1.98	1.92
N = 1000000	19.9	20.78

Other Sequence Containers

The STL also includes priority queue, queue, and linked list classes, but those aren't too important to us right now.
